

Scikit Learn Notes

Stephen CUI

October, 25, 2022

Contents

I	User Guide	1
1	Supervised learning	3
1.1	Linear Models	3
1.1.1	Ordinary Least Squares	3
1.1.2	Ridge regression and classification	4
1.1.3	Lasso	4
1.1.4	Multi-task Lasso	4
1.1.5	Logistic regression	4
2	Model selection and evaluation	5
2.1	Cross-validation: evaluating estimator performance	5
2.2	Tuning the hyper-parameters of an estimator	5
II	Example	7
3	Generalized Linear Models	9
3.1	Linear Regression Example	9
3.2	Non-negative least squares	9

Part I

User Guide

Chapter 1

Supervised learning

1.1 Linear Models

The following are a set of methods intended for regression in which the target value is expected to be a linear combination of the features. In mathematical notation, if \hat{y} is the predicted value.

$$\hat{y}(w, x) = w_0 + w_1x_1 + \dots + w_px_p \quad (1.1)$$

Across the module, we designate the vector $w = (w_1, \dots, w_p)$ as `coef_` and w_0 as `intercept_`.

To perform classification with generalized linear models, see [Logistic regression](#).

1.1.1 Ordinary Least Squares

LinearRegression fits a linear model with coefficients $w = (w_1, \dots, w_p)$ to minimize the residual sum of squares between the observed targets in the dataset, and the targets predicted by the linear approximation. Mathematically it solves a problem of the form:

$$\min_w \|Xw - y\|_2^2 \quad (1.2)$$

LinearRegression will take in its `fit` method arrays `X`, `y` and will store the coefficients w of the linear model in its `coef_` member:

```
from sklearn import linear_model
reg = linear_model.LinearRegression()
reg.fit([[0, 0], [1, 1], [2, 2]], [0, 1, 2])
reg.coef_
```

The coefficient estimates for Ordinary Least Squares **rely on the independence of the features**. When features are correlated and the columns of the design matrix X have an approximately linear dependence, the design matrix becomes close to singular and as a result, the least-squares estimate becomes highly sensitive to random errors in the observed target, producing a large variance. This situation of multicollinearity can arise, for example, when data are collected without an experimental design.

Examples:

- [Linear Regression Example](#)

Non-Negative Least Squares

It is possible to constrain all the coefficients to be non-negative, which may be useful when they represent some physical or naturally non-negative quantities (e.g., frequency counts or prices of goods). LinearRegression accepts a boolean `positive` parameter: when set to `True` Non-Negative Least Squares are then applied. **Examples:**

- [Non-negative least squares](#)

Ordinary Least Squares Complexity**1.1.2 Ridge regression and classification****1.1.3 Lasso****1.1.4 Multi-task Lasso****1.1.5 Logistic regression**

Chapter 2

Model selection and evaluation

2.1 Cross-validation: evaluating estimator performance

Learning the parameters of a prediction function and testing it on the same data is a methodological mistake: a model that would just repeat the labels of the samples that it has just seen would have a perfect score but would fail to predict anything useful on yet-unseen data. This situation is called **overfitting**. To avoid it, it is common practice when performing a (supervised) machine learning experiment to hold out part of the available data as a **test set** $X_{\text{test}}, y_{\text{test}}$. Note that the word “experiment” is not intended to denote academic use only, because even in commercial settings machine learning usually starts out experimentally. Here is a flowchart of typical cross validation workflow in model training. The best parameters can be determined by grid search techniques in [Tuning the hyper-parameters of an estimator](#).

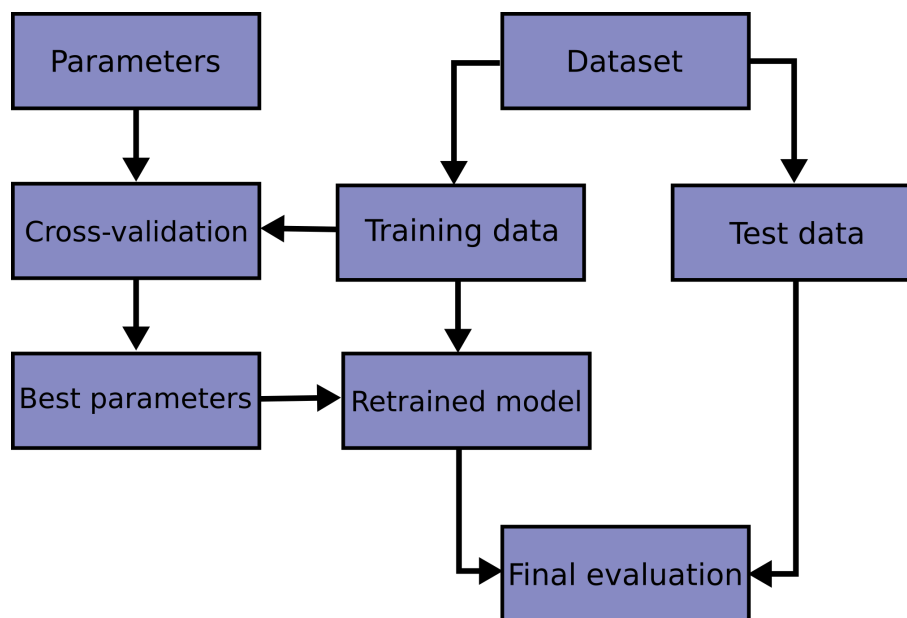


Figure 2.1: https://scikit-learn.org/stable/_images/grid_search_workflow.png

2.2 Tuning the hyper-parameters of an estimator

Part II

Example

Chapter 3

Generalized Linear Models

3.1 Linear Regression Example

The example below uses only the first feature of the `diabetes` dataset, in order to illustrate the data points within the two-dimensional plot. The straight line can be seen in the plot, showing how linear regression attempts to draw a straight line that will best minimize the residual sum of squares between the observed responses in the dataset, and the responses predicted by the linear approximation.

The coefficients, residual sum of squares and the coefficient of determination are also calculated.

3.2 Non-negative least squares

In this example, we fit a linear model with positive constraints on the regression coefficients and compare the estimated coefficients to a classic linear regression.